



# 703308 VO High-Performance Computing WS2022/2023

## MPI - Message Passing Interface


Philipp Gschwandtner

# Overview

---

- ▶ general concepts about MPI
  - ▶ characteristics
  - ▶ programming model
  - ▶ startup
- ▶ point-to-point communication
- ▶ collective communication
- ▶ practical example

# Motivation for using MPI: data distribution



# Message Passing Interface (MPI)

---

- ▶ message passing library for distributed memory parallelism
- ▶ de-facto standard for C/C++ and Fortran
- ▶ maintained by the MPI Forum
  - ▶ initial release in 1994 (version 1.0)
  - ▶ updates in 1997 (2.0), 2012 (3.0), 2015 (3.1) and 202x (4.0)
  - ▶ specification updates slow, aim at stability and high TRL
    - ▶ “On June 30 2020, the MPI forum voted for version 3.3 of [these voting rules](#) (effective June 30th, 2020).”

# MPI implementations

---

- ▶ **OpenMPI**

- ▶ open source
- ▶ merge of multiple previous MPI implementations
- ▶ default on many systems

- ▶ **MPICH**

- ▶ also open source
- ▶ basis for many vendor implementations such as Intel, IBM, Cray, Microsoft, ...
  - ▶ default on many systems

- ▶ Do not confuse implementation versions with specification versions!
- ▶ Do not confuse implementation adherence with specification adherence!

# Main characteristics

---

- ▶ offers specific tools for
  - ▶ sending and receiving messages
  - ▶ waiting and synchronization
  - ▶ identification of individual processes
  - ▶ ...
- ▶ additional convenience tools for
  - ▶ partitioning and distributing data
  - ▶ organizing processes in structures
  - ▶ large-scale I/O operations
  - ▶ ...

## Main characteristics cont'd

---

- ▶ a lot of user responsibility

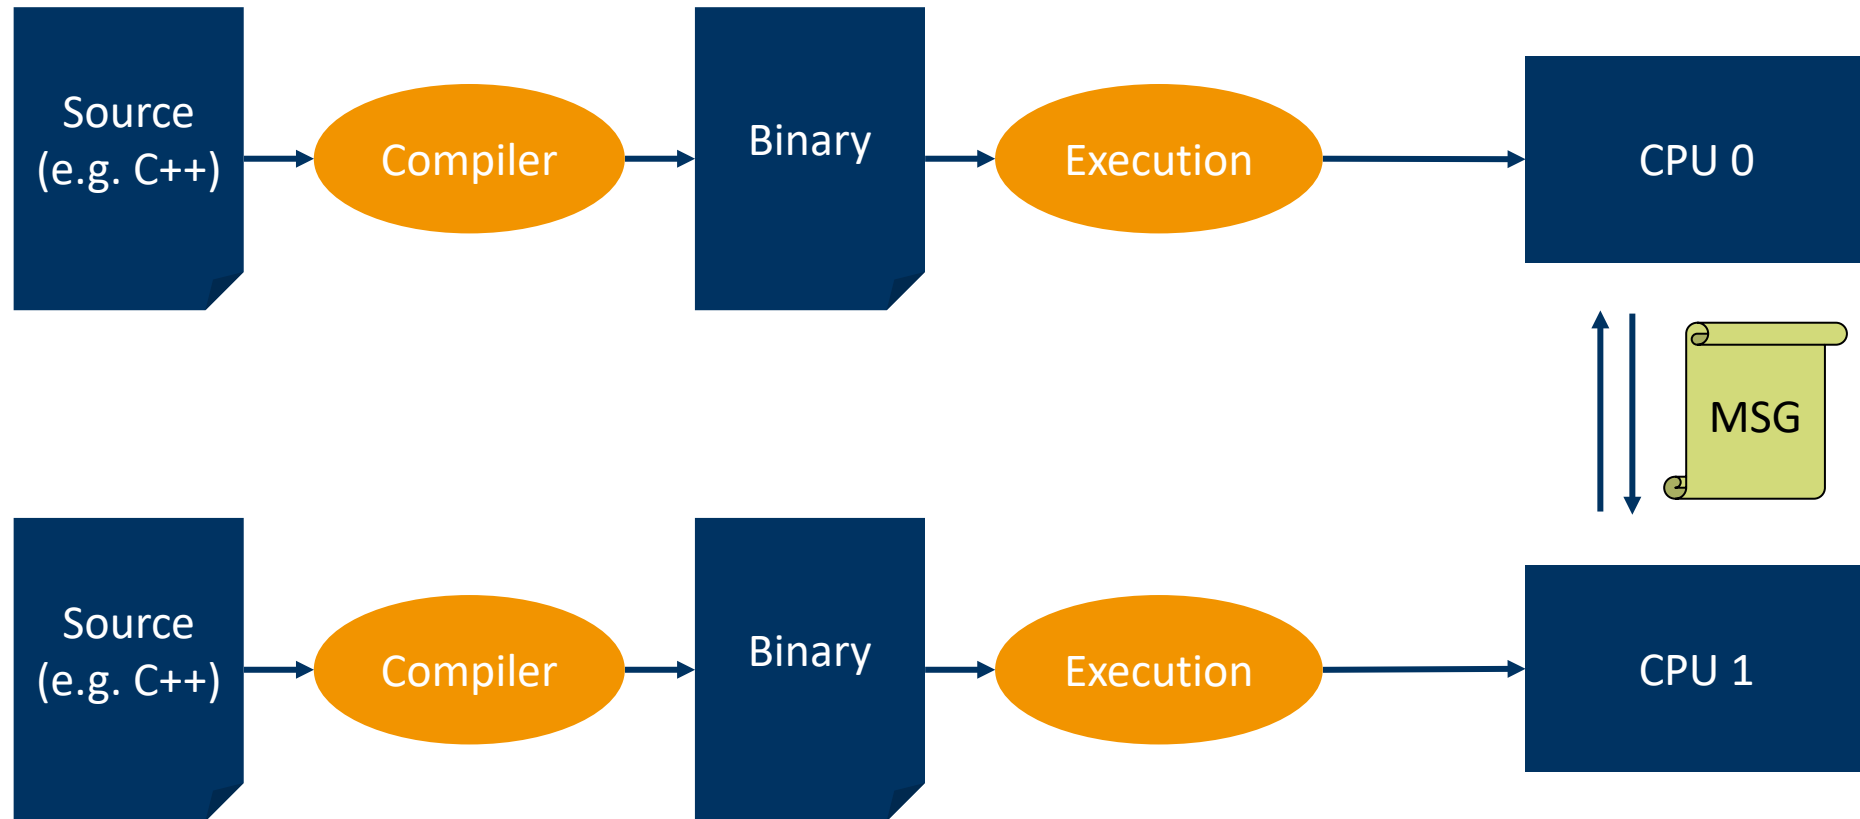
- ▶ explicit parallelism and communication
- ▶ program correctness
- ▶ performance optimization
  - ▶ (non-)blocking
  - ▶ (a)synchronous

- ▶ a lot of advantages

- ▶ available everywhere
- ▶ several implementations
- ▶ portable to many architectures
- ▶ very high performance

## Recap: MIMD: MPMD


---





## Recap: MIMD: SPMD

---



## MPMD through SPMD

---

- ▶ many MPI implementations support only SPMD
- ▶ SPMD can emulate MPMD

```
int main() {  
    // get id information  
    int cpuID = ...;  
    if(cpuID==0) {  
        ... // program A  
    } else {  
        ... // program B  
    }  
}
```

# Parallelism requires two mechanisms

---

- ▶ a mechanism for spawning processes
  - ▶ multiple ways of achieving this
  - ▶ we won't look at this in too much detail
  - ▶ simply rely on `mpiexec` to do the work for you
- ▶ a mechanism for sending and receiving messages
  - ▶ many, many ways of exchanging messages
  - ▶ we will definitely look at this in a lot of detail
  - ▶ tons of functionality to choose from, as we'll see in a bit

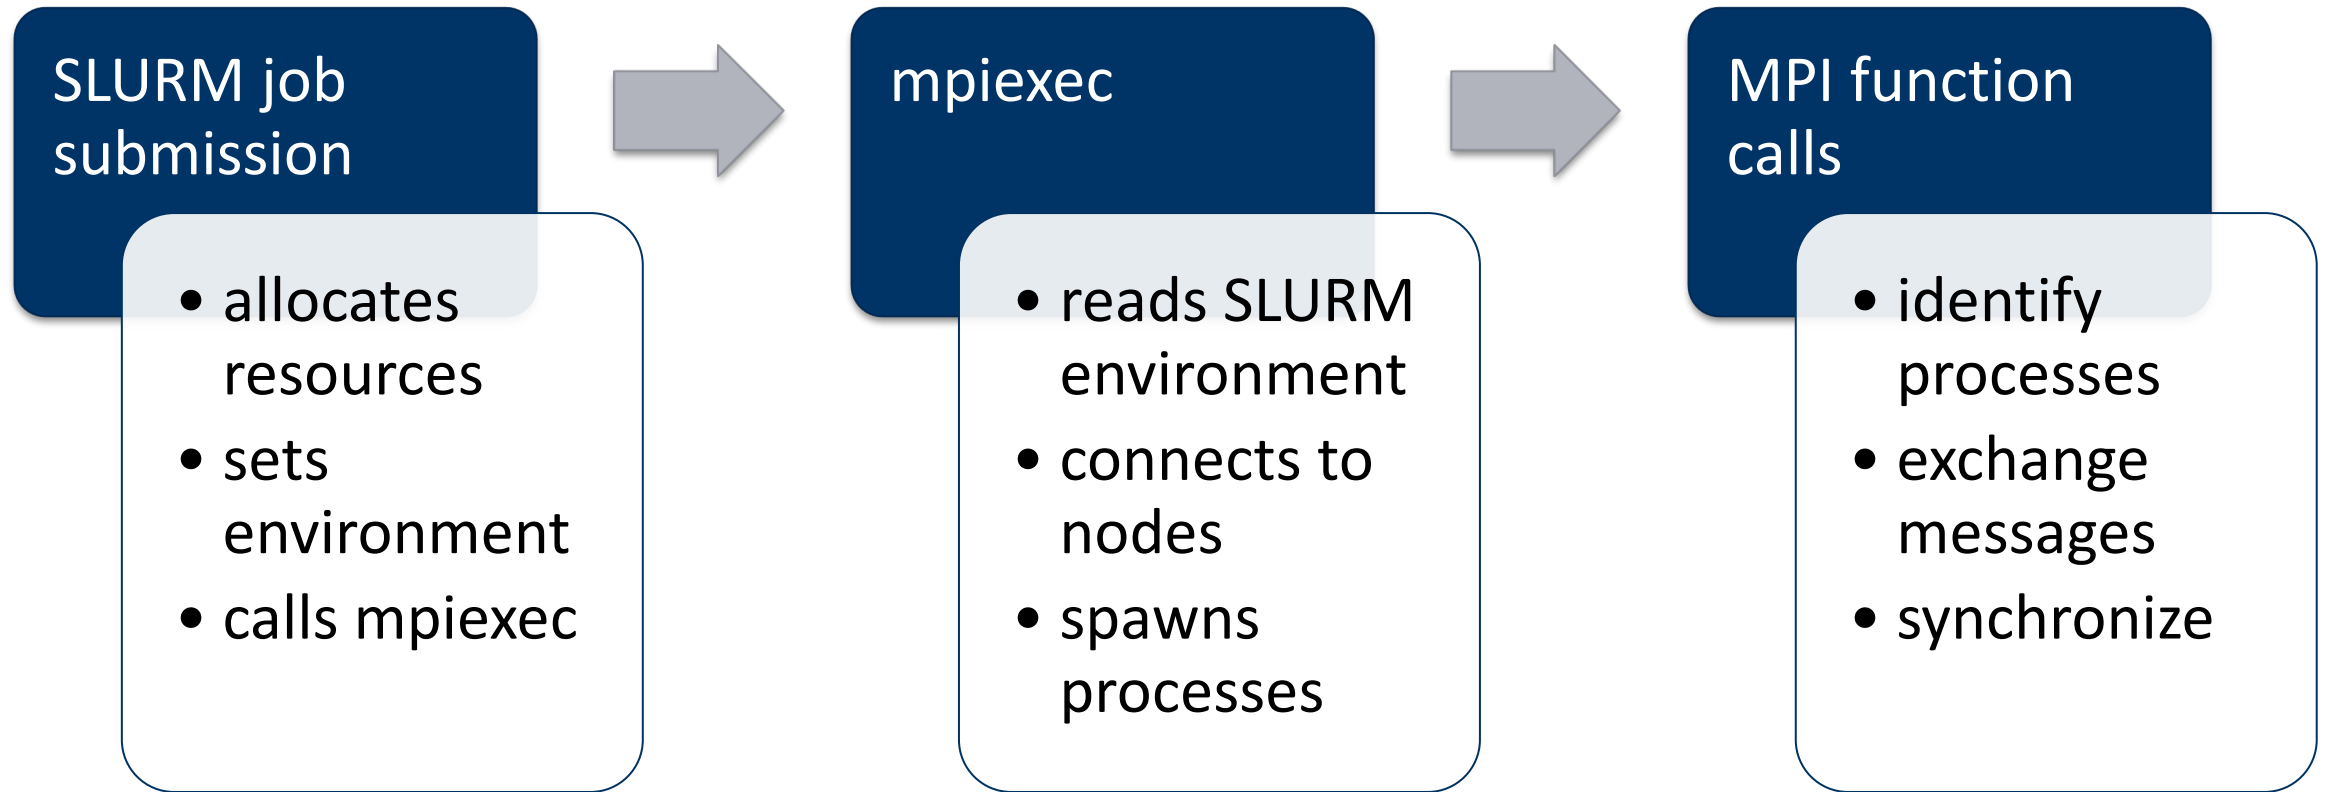
# Compiler and execution wrappers

---

- ▶ `mpicc / mpic++` for compiling
  - ▶ OpenMPI: `--showme` prints compiler flags
  - ▶ passes all additional compiler flags to backend compiler (e.g. `mpicc -g`)
- ▶ `mpiexec` for running
  - ▶ formerly `mpirun`, but `mpiexec` is standardized

# Startup procedure of an MPI application

---



# Hello world in MPI

---

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {
    MPI_Init(&argc, &argv); // initialize the MPI environment
    int size;
    MPI_Comm_size(MPI_COMM_WORLD, &size); // get the number of ranks
    int rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank); // get the rank of the caller
    printf("Hello world from rank %d of %d\n", rank, size);
    MPI_Finalize(); // cleanup
}
```

## Setup and teardown

---

- ▶ `int MPI_Init( int* argc, char*** argv )`
  - ▶ must be called by every process before calling any other MPI function
  - ▶ initializes the MPI library
- ▶ `int MPI_Finalize( void )`
  - ▶ must be the last MPI function called by every process
  - ▶ user must ensure completion of all locally (!) pending communication
  - ▶ performs library cleanup

# Who am I talking to?

---

- ▶ in MPI-speak, processes are known as “*ranks*”
  - ▶ numbered from 0 to N-1
  - ▶ own rank can be queried with

```
int MPI_Comm_rank( MPI_Comm comm, int* rank )
```
  - ▶ number of ranks can be queried with

```
int MPI_Comm_size( MPI_Comm comm, int* size)
```
- ▶ almost all MPI semantics are relative to a “*communicator*” or “*group*”
  - ▶ identifies a set of ranks
  - ▶ always available: MPI\_COMM\_WORLD (=everyone)
  - ▶ new communicators and groups that hold subsets of ranks can be created
  - ▶ when developing a library, always create your own communicator!
    - ▶ frequent MPI programming pattern






# Point-to-Point Communication



# Point-to-point communication

---

- ▶ `MPI_Send(...)/MPI_Recv(...)`
  - ▶ single sender, single receiver  
(*“point-to-point”*)
- ▶ simplest form of communication available
  - ▶ not necessarily the best
- ▶ multiple different types
  - ▶ (a)synchronous
  - ▶ (non-)blocking



## MPI\_Send

---

- ▶ `int MPI_Send(const void* buf, int count, MPI_Datatype datatype, int dest, int tag, MPI_Comm comm)`
  - ▶ `buf`: source buffer to send data from
  - ▶ `count`: number of data elements to send
  - ▶ `datatype`: type of data to send
  - ▶ `dest`: destination rank
  - ▶ `tag`: user-defined message type or category
  - ▶ `comm`: communicator
  
- ▶ `MPI_Send(&number, 1, MPI_INT, 1, 42, MPI_COMM_WORLD);`

## MPI\_Recv

---

- ▶ `int MPI_Recv(void* buf, int count, MPI_Datatype datatype, int source, int tag, MPI_Comm comm, MPI_Status* status)`
  - ▶ `buf`: destination buffer to save data to
  - ▶ `count`: number of data elements to receive
  - ▶ `datatype`: type of data to receive
  - ▶ `source`: source rank
  - ▶ `tag`: user-defined message type or category
  - ▶ `comm`: communicator
  - ▶ `status`: holds additional information (e.g. rank of sender or tag of message)
- ▶ `MPI_Recv(&number, 1, MPI_INT, 0, 42, MPI_COMM_WORLD, MPI_STATUS_IGNORE);`

## Basic send/receive example

---

```
int number;
if (rank == 0) {
    number = -1;
    MPI_Send(&number, 1, MPI_INT, 1, 42, MPI_COMM_WORLD);
} else if (rank == 1) {
    MPI_Recv(&number, 1, MPI_INT, 0, 42, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    printf("Rank 1: Received %d from rank 0\n", number);
}
```

## Side note: MPI data type conversion

---

- ▶ “representation conversion”
  - ▶ changes the binary representation of some data
  - ▶ c.f. C++ reinterpret casts
  - ▶ allowed in MPI, required by supporting heterogeneous architectures
    - ▶ e.g. sender is ARM, receiver is x86, ...
    - ▶ e.g. big-endian vs. little-endian, ASCII vs. EBCDIC, ...
- ▶ “type conversion”
  - ▶ converts the actual data type, e.g. from float to integer
  - ▶ c.f. C++ numeric casts
  - ▶ not allowed in MPI

# Predefined MPI constants

---

- ▶ datatypes

- ▶ MPI\_INT, MPI\_FLOAT, MPI\_DOUBLE, MPI\_BYTE, MPI\_CHAR, ...

- ▶ wildcards & misc

- ▶ MPI\_ANY\_SOURCE
- ▶ MPI\_ANY\_TAG
- ▶ MPI\_COMM\_WORLD
- ▶ MPI\_STATUS\_IGNORE
- ▶ ...

# (Non-)Blocking and (a)synchronous communication


---

- ▶ distinguish two important properties
  - ▶ When does the MPI function call return?
    - ▶ Can I overwrite the send buffer?
    - ▶ When is all the data in the receive buffer?
  - ▶ When does the corresponding message transfer happen?
    - ▶ Do I need to wait for the receiver to get the entire message?
    - ▶ Do I need to wait for the receiver to begin receiving?

```
if (rank == 0) {  
    MPI_Send(&number, ...);  
} else if (rank == 1) {  
    MPI_Recv(&number, ...);  
}
```



## (Non-)Blocking and (a)synchronous communication cont'd



# Blocking vs. non-blocking communication

---

- ▶ blocking point-to-point:  
MPI\_Send() and MPI\_Recv()
  - ▶ allows to re-use send buffer after send call returns
  - ▶ allows to read receive buffer after receive call returns

```
if (rank == 0) {  
    MPI_Send(&number, ...);  
    // re-use number here  
} else if (rank == 1) {  
    MPI_Recv(&number, ...);  
    // use number here  
}
```

## Blocking vs. non-blocking communication cont'd

---

- ▶ non-blocking point-to-point:  
MPI\_Isend() and MPI\_Irecv()
  - ▶ send and receive return (almost) immediately
  - ▶ MPI\_Wait() call blocks until buffer can be read/re-used

```
MPI_Request request;
if (rank == 0) {
    MPI_Isend(&number, ..., &request);
    MPI_Wait(&request, MPI_STATUS_IGNORE);
    // re-use number here
} else if (rank == 1) {
    MPI_Irecv(&number, ..., &request);
    MPI_Wait(&request, MPI_STATUS_IGNORE);
    // re-use number here
}
```

## (A)Synchronous send modes


---

- ▶ `MPI_Ssend()` – synchronous mode
    - ▶ will wait for matching receive
  - ▶ `MPI_Bsend()` – buffered mode
    - ▶ will buffer, won't wait for a matching receive
  - ▶ `MPI_Rsend()` – ready mode
    - ▶ requires an already posted, matching receive (developer responsibility!)
  - ▶ `MPI_Send()` – standard mode
    - ▶ may buffer (depends on message size)
    - ▶ may or may not wait for matching receive
- ▶ and there are also non-blocking variants for ALL of them...

# Message order preservation

---

- ▶ messages do NOT overtake each other if
  - ▶ same communicator
  - ▶ same source rank
  - ▶ same destination rank
- ▶ regardless of blocking or synchronization mode
- ▶ mandated by MPI specification





# Collective Communication




# Collective communication

---

- ▶ convenience function for frequently-used programming patterns (e.g. distributing data)
  - ▶ can involve several ranks at the same time, not just 2
  - ▶ must be called by ALL ranks in the communicator
  - ▶ must be called in-order by all ranks (no interleaving of multiple collective communication calls!)
  - ▶ is locally finished when local operation has finished
  - ▶ is globally finished when all participating ranks are finished
  - ▶ available as blocking and non-blocking variants (but cannot be mixed!)

## MPI\_Scatter/MPI\_Scatterv

- ▶ sends chunks of data to multiple ranks, including root itself
- ▶ simple way of partitioning and distributing data
- ▶ MPI\_Scatterv() allows varying counts of elements to be distributed to each rank





# MPI\_Scatter

---

- ▶ `int MPI_Scatter(const void* sendbuf, int sendcount, MPI_Datatype sendtype, void* recvbuf, int recvcount, MPI_Datatype recvtype, int root, MPI_Comm comm)`
  - ▶ `sendbuf`: source buffer to send data from
  - ▶ `sendcount`: number of data elements to send to each rank
  - ▶ `sendtype`: type of data to send
  - ▶ `recvbuf`: destination buffer to save data to
  - ▶ `recvcount`: number of data elements to receive at each rank
  - ▶ `recvtype`: type of data to receive
  - ▶ `root`: rank of the sender
  - ▶ `comm`: communicator

## Scatter example

---


```
int globaldata[4];

if(rank==0) {
    for(int i = 0; i < 4; i++) {
        globaldata[i] = ...
    }
}

int localdata;
MPI_Scatter(globaldata, 1, MPI_INT, &localdata, 1, MPI_INT, 0,
            MPI_COMM_WORLD);
```

# MPI\_Gather/MPI\_Gatherv


- ▶ sends chunks of data from multiple ranks, including root itself, to root
- ▶ simple way of collecting data
- ▶ MPI\_Gatherv() allows varying counts of elements to be collected from each rank



# MPI\_Bcast

---

- ▶ broadcast operation
- ▶ sends copies of data to multiple ranks



## Q: emulating broadcast with point-to-point?


---

```
MPI_Bcast(&buf, 1, MPI_INT, 0, MPI_COMM_WORLD);  
// ##### or instead: #####  
if (rank == 0) {  
    for (int i = 0; i < size; i++) {  
        if (i != rank) {  
            MPI_Send(&buf, 1, MPI_INT, i, 0, MPI_COMM_WORLD);  
        }  
    }  
} else {  
    MPI_Recv(&buf, 1, MPI_INT, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
}
```


# Collective communication patterns

---


- ▶ chosen automatically at runtime by MPI implementation
- ▶ may depend on multiple parameters, including
  - ▶ type of operation (e.g. broadcast)
  - ▶ number and location of ranks
  - ▶ size and structure of data
  - ▶ hardware capabilities



sequential algorithm  
 $O(\text{num\_ranks})$



tree-based algorithm  
e.g.  $O(\log_2(\text{num\_ranks}))$



hardware operation  
 $O(1)$


# Barrier

---

- ▶ `int MPI_Barrier(MPI_Comm comm)`
  - ▶ `comm`: communicator
- ▶ causes all ranks to wait until everyone reached the barrier
  - ▶ normally not needed: explicit data communication usually also inherently synchronizes
  - ▶ often used for debugging and profiling
    - ▶ Don't forget to remove in production/release!

# MPI\_Reduce

- ▶ aggregate data from multiple ranks, including root itself, to root
  - ▶ e.g. MPI\_SUM
  - ▶ may use optimized collective communication patterns
- ▶ requires associative reduction operation ◦ such that e.g.
$$(x_0 \circ x_1) \circ (x_2 \circ x_3) = ((x_0 \circ x_1) \circ x_2) \circ x_3$$
- ▶ be careful with floating point types!





## MPI\_Reduce cont'd

---

- ▶ `int MPI_Reduce(const void* sendbuf, void* recvbuf, int count, MPI_Datatype datatype, MPI_Op op, int root, MPI_Comm comm)`
  - ▶ `sendbuf`: source buffer to reduce data from
  - ▶ `recvbuf`: destination buffer to reduce data into
  - ▶ `count`: number of data elements in source and destination buffers
  - ▶ `datatype`: type of data to reduce
  - ▶ `op`: reduction operation
  - ▶ `root`: rank of the destination of aggregated result
  - ▶ `comm`: communicator

# Available reduction operations

---


- ▶ several pre-defined
  - ▶ MPI\_MAX, MPI\_MIN
  - ▶ MPI\_SUM, MPI\_PROD
  - ▶ MPI\_LAND, MPI\_LOR, MPI\_LXOR
  - ▶ MPI\_BAND, MPI\_BOR, MPI\_BXOR
  - ▶ MPI\_MAXLOC, MPI\_MINLOC
- ▶ All pre-defined operations are associative and commutative
- ▶ also user-defined ops are possible
  - ▶ must be associative
  - ▶ may be commutative
  - ▶ requires a specific function signature
  - ▶ needs to be registered as an MPI handle

# Additional MPI functions

---

- ▶ **MPI\_Wtime**
  - ▶ returns time in seconds since an arbitrary time in the past (Wall clock time)
- ▶ **MPI\_Sendrecv**
  - ▶ convenience wrapper for blocking send and receive
- ▶ **MPI\_Allreduce/MPI\_Allgather/...**
  - ▶ same as non-all versions, but result is available everywhere (performance impact!)
- ▶ **MPI\_Scan/MPI\_Exscan**
  - ▶ inclusive and exclusive prefix reductions
- ▶ **MPI\_Wait/MPI\_Test**
  - ▶ blocking/non-blocking check whether pending operation completed
- ▶ **MPI\_Probe/MPI\_Iprobe**
  - ▶ blocking/non-blocking check for new message without actually receiving it

## Example code: naïve matrix multiplication



```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>
```

```
#define SIZE 4
```

```
int A[SIZE][SIZE];
int B[SIZE][SIZE];
int C[SIZE][SIZE];
```

```
void fill_matrix(int m[SIZE][SIZE]);
void print_matrix(int m[SIZE][SIZE]);
```

## Example code: naïve matrix multiplication cont'd

---

```
int myRank, numProcs;

MPI_Init(&argc, &argv);
MPI_Comm_rank(MPI_COMM_WORLD, &myRank);
MPI_Comm_size(MPI_COMM_WORLD, &numProcs);
// if matrix size not divisible
if(SIZE % numProcs != 0) {
    MPI_Finalize();
    return EXIT_FAILURE;
}
// root generates input data
if(myRank == 0) {
    fill_matrix(A);
    fill_matrix(B);
}
```

```
// compute boundaries of local computation
int start = myRank*SIZE/numProcs;
int end = (myRank+1)*SIZE/numProcs;

// distribute rows of A to everyone
MPI_Scatter(A, SIZE*SIZE/numProcs, MPI_INT,
    A[start], SIZE*SIZE/numProcs, MPI_INT, 0,
    MPI_COMM_WORLD);

// send entire matrix B to everyone
MPI_Bcast(B, SIZE*SIZE, MPI_INT, 0,
    MPI_COMM_WORLD);
```

## Example code: naïve matrix multiplication cont'd

---

```
// local computation of every rank
for(int i = start; i < end; i++) {
    for(int j = 0; j < SIZE; j++) {
        C[i][j] = 0;
        for(int k = 0; k < SIZE; k++) {
            C[i][j] += A[i][k] * B[k][j];
        }
    }
}
```

```
// gather result rows back to root
MPI_Gather(C[start], SIZE*SIZE/numProcs,
          MPI_INT, C, SIZE*SIZE/numProcs, MPI_INT, 0,
          MPI_COMM_WORLD);

if(myRank == 0) { print_matrix(C); }

MPI_Finalize();

return EXIT_SUCCESS;
```

# Submitting to a cluster (SLURM & SGE)

---

```
#!/bin/bash

# submission partition
#SBATCH --partition std.q
# name of the job
#SBATCH --job-name my_test_job
# redirect output stream to this file
#SBATCH --output output.dat
# specify parallel environment
#SBATCH --ntasks 8
#SBATCH --ntasks-per-node 2

srun /path/to/application
# or
mpiexec -n $SLURM_NTASKS /path/to/application
```

```
#!/bin/bash

# submission queue
#$ -q std.q
# change to current directory
#$ -cwd
# name of the job
#$ -N my_test_job
# redirect output stream to this file
#$ -o output.dat
# join the output and error stream
#$ -j yes
# specify parallel environment
#$ -pe openmpi-2perhost 8

mpiexec -n 8 /path/to/application
```

# Summary

---

- ▶ general concepts about MPI
  - ▶ characteristics
  - ▶ program model
  - ▶ startup
- ▶ point-to-point communication
- ▶ collective communication
- ▶ practical example (matrix multiplication)